

Introduction to Embedded Systems

4. Programming Paradigms

Prof. Dr. Marco Zimmerling



Organization

Join ILIAS course:

- Login: RZ username + password
- Course password: **es-0x8af**

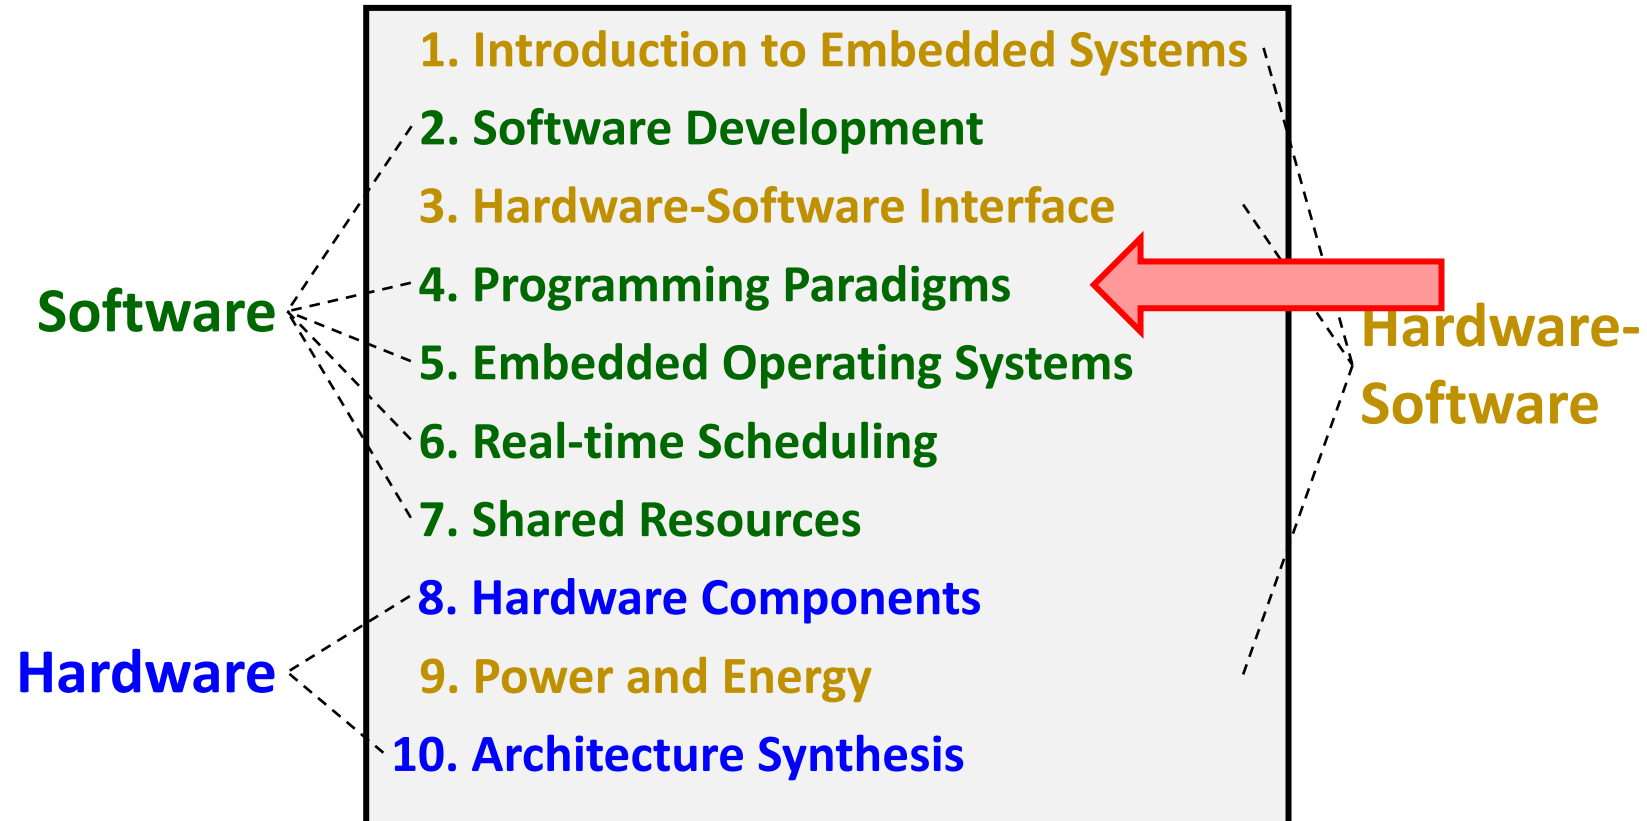
Exercises:

- From 12:00 to 14:00 in [HS 00-036 \(English\)](#) and [SR 02-016/18 \(German\)](#)
- Today (November 22):
 - Solutions of second exercise sheet presented
- Next week (November 29):
 - Third exercise sheet released + overview of tasks given



Exam: [March 4, 2023](#) from 10:00 to 12:00, five rooms in Georges-Köhler-Allee 101


Where we are ...




Reactive Systems and Timing

Real-Time Systems

In many *cyber-physical systems (CPSs)*, correct timing is a matter of *correctness*, not performance: *an answer arriving too late is consider to be an error*.



Methods to Determine the Execution Time of a Task



Different Programming Paradigms


Overview

- There are many *structured ways of programming an embedded system*.
- In this lecture, only the main principles will be covered:
 - *time triggered approaches*
 - periodic
 - cyclic executive
 - generic time-triggered scheduler
 - *event triggered approaches*
 - non-preemptive
 - preemptive – stack policy
 - preemptive – cooperative scheduling
 - preemptive – multitasking

Time-Triggered Systems


Pure time-triggered model:

- *no interrupts* are allowed, except by timers
- the *schedule* of tasks is *computed off-line* and therefore, complex sophisticated algorithms can be used
- the scheduling at run-time is fixed and therefore, it is *deterministic*
- the interaction with environment happens through *polling*



Simple Periodic TT Scheduler

- A *timer interrupts regularly* with period P .
- All tasks have *same period P* .




- *Properties:*
 - later tasks, for example T_2 and T_3 , have unpredictable starting times
 - the communication between tasks or the use of common resources is safe, as there is a static ordering of tasks, for example T_2 starts after finishing T_1
 - as a necessary precondition, the sum of WCETs of all tasks within a period is bounded by the period P :

$$\sum_{(k)} WCET(T_k) < P$$

Time-Triggered Cyclic Executive Scheduler


- Suppose now that *tasks may have different periods*.
- To accommodate this situation, the *period P is partitioned into frames of length f* .




- We have a *problem* to determine a feasible schedule if there are *tasks with a long execution time*.
 - long tasks could be partitioned into a sequence of short sub-tasks
 - but this is a tedious and error-prone process, as the local state of the task must be extracted and stored globally

Time-Triggered Cyclic Executive Scheduling

- *Example* of a single periodic task τ_i :



- *A set of periodic tasks* Γ :



Time-Triggered Cyclic Executive Scheduling

Some conditions for period P and frame length f :

- A task executes at most once within a frame:

$$f \leq T_i \quad \forall \text{ tasks } \tau_i$$

period of task

- P is a multiple of f .

- Period P is least common multiple of all periods T_k .

- Tasks start and complete within a single frame:

$$f \geq C_i \quad \forall \text{ tasks } \tau_i$$


worst case execution time
of task

- Between release time and deadline of every task there is at least one full frame:

$$2f - \gcd(T_i, f) \leq D_i \quad \forall \text{ tasks } \tau_i$$

relative deadline of task

Sketch of Proof for Last Condition



Example: Cyclic Executive Scheduling

Conditions:

$$f \leq \min\{4, 5, 20\} = 4$$


$$f \geq \max\{1.0, 1.0, 1.8, 2.0\} = 2.0$$

$$2f - \gcd(T_i, f) \leq D_i \quad \forall \text{ tasks } \tau_i$$

possible solution: $f = 2$

Γ	T_i	D_i	C_i
τ_1	4	4	1.0
τ_2	5	5	1.8
τ_3	20	20	1.0
τ_4	20	20	2.0

Feasible solution ($f=2$):



Time-Triggered Cyclic Executive Scheduling

Checking for correctness of schedule:

- f_{ij} denotes the number of the frame in which that instance j of task τ_i executes.
- Is P a common multiple of all periods T_i ?
- Is P a multiple of f ?
- Is the frame sufficiently long?

$$\sum_{\{i \mid f_{ij}=k\}} C_i \leq f \quad \forall 1 \leq k \leq \frac{P}{f}$$

- Determine initial phases such that instances of tasks start after their release time:


$$\Phi_i = \min_{1 \leq j \leq P/T_i} \{(f_{ij} - 1)f - (j - 1)T_i\} \quad \forall \text{ tasks } \tau_i$$

- Are deadlines respected?

$$(j - 1)T_i + \Phi_i + D_i \geq f_{ij}f \quad \forall \text{ tasks } \tau_i, 1 \leq j \leq P/T_i$$

Generic Time-Triggered Scheduler

- In an *entirely time-triggered system*, the temporal control structure of all tasks is established a priori by off-line support-tools.
- This *temporal control structure is encoded in a Task-Descriptor List (TDL)* that contains the cyclic schedule for all activities of the node.
- This *schedule* considers the required precedence and mutual exclusion relationships among the tasks such that an explicit coordination of the tasks by the operating system at run time is not necessary.
- *The dispatcher is activated by a synchronized clock tick.* It looks at the TDL, and then performs the action that has been planned for this instant [Kopetz].



start times can be arbitrarily chosen within the period

Simplified Time-Triggered Scheduler

main:

```
determine static schedule  $(t(k), T(k))$ , for  $k=0,1,\dots,n-1$ ;  
determine period of the schedule  $P$ ;  
set  $i=k=0$  initially; set the timer to expire at  $t(0)$ ;  
while (true) sleep();
```

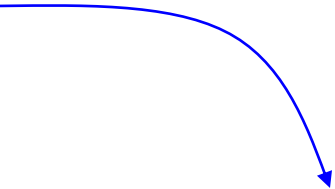
usually done offline

Timer Interrupt:

```
 $k_{old} := k$ ;  
 $i := i+1$ ;  $k := i \bmod n$ ;  
set the timer to expire at  $\lfloor i/n \rfloor * P + t(k)$ ;  
execute task  $T(k_{old})$ ;  
return;
```

set CPU to low power mode;
processing continues after interrupt

for example using a function pointer in C;
task returns after finishing.



k	$t(k)$	$T(k)$
0	0	T_1
1	3	T_2
2	7	T_1
3	8	T_3
4	12	T_2

$n=5, P = 16$

Summary Time-Triggered Scheduler

Advantages:

- *deterministic schedule*: conceptually simple (static table); easy to validate, test, and certify
- *no problems* in using *shared resources*

Disadvantages:

- external communication only via *polling*
- *inflexible* as no adaptation to the environment
- serious *problems* if there are *long tasks*


Extensions:

- *allow interrupts* → be careful with shared resources and the WCET of tasks!!
- *allow preemptable* background tasks
- *check for task overruns* (execution time longer than WCET) using a watchdog timer

Event-Triggered Systems

The schedule of tasks is determined by the occurrence of external or internal events:

- *dynamic and adaptive*: there are possible problems with respect to timing, the use of shared resources and buffer over- or underflow
- *guarantees* can be given either off-line (if bounds on the behavior of the environment are known) or during run-time



Non-Preemptive Event-Triggered Scheduling

Principle:

- To each event, there is associated a corresponding task that will be executed.
- Events are emitted by (a) external interrupts or (b) by tasks themselves.
- All events are collected in a single queue.
- Depending on the queuing discipline (e.g., first come first serve), an event is chosen for execution, that is, the corresponding task is executed.
- A running task *cannot be preempted* by another task. It can only be preempted by an interrupt that registers an event and puts it into the queue.

Extensions:

- A *background task*, which has the lowest priority, can run if the event queue is empty. It will be preempted by any event processing.
- *Timed events* are ready for execution only after a certain time interval has elapsed. This enables, for example, periodic task instantiations.

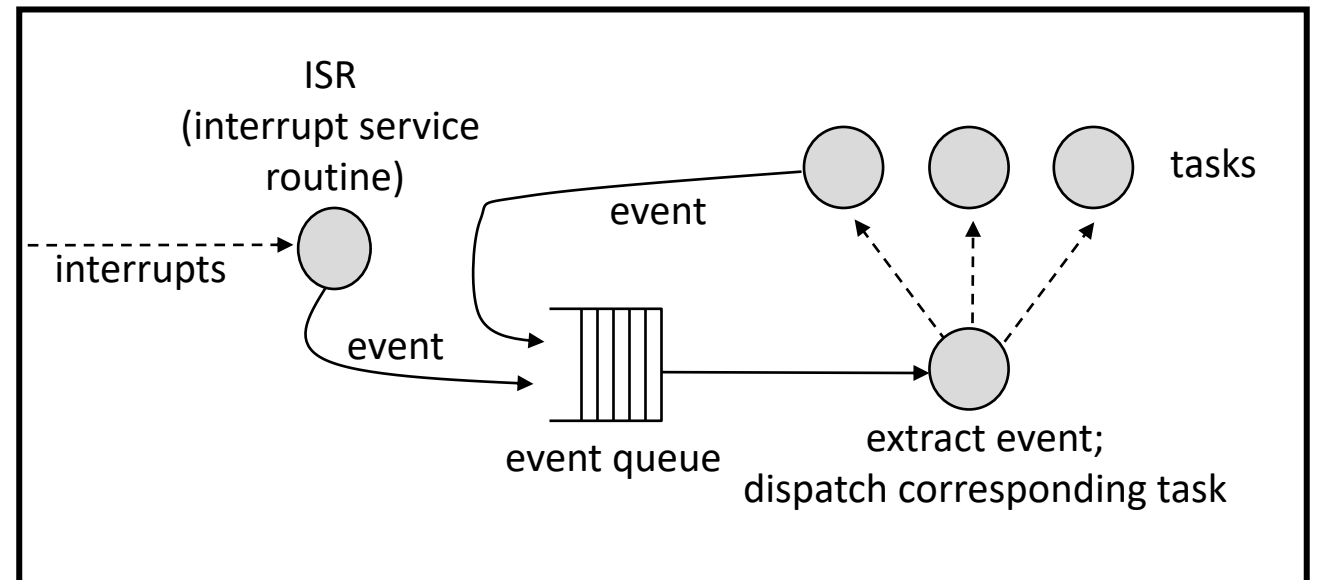
Non-Preemptive Event-Triggered Scheduling

```
main:
  while (true) {
    if (event queue is empty) {
      sleep();
    } else {
      extract event from event queue;
      execute task corresponding to event;
    }
  }
}
```

set the CPU to low power mode;
continue processing after interrupt

for example using a function pointer in C;
task returns after finishing.


```
Interrupt:
  put event into event queue;
  return;
```



Non-Preemptive Event-Triggered Scheduling

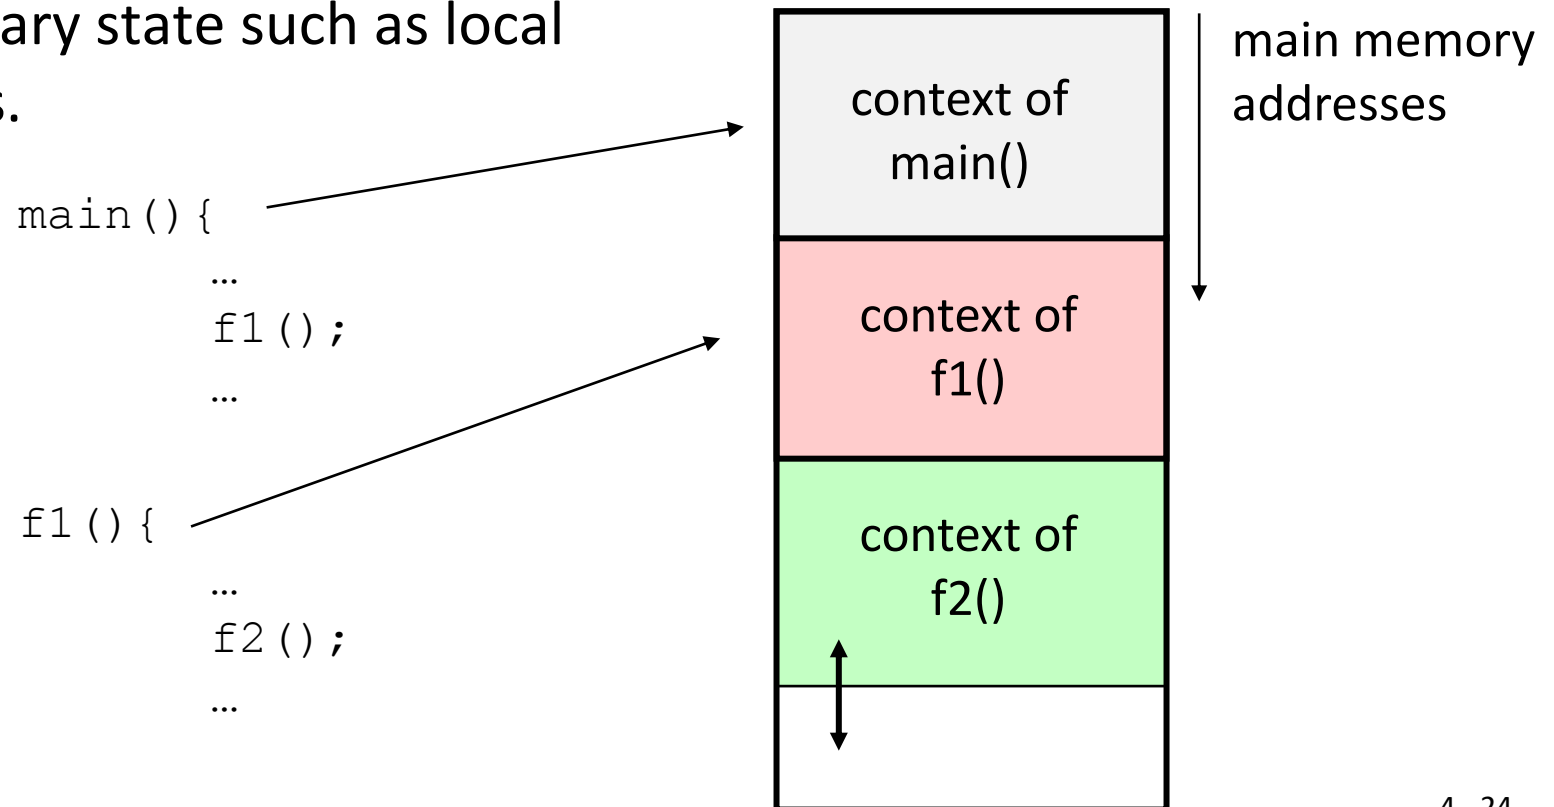
Properties:

- *communication between tasks* does not lead to a simultaneous access to shared resources, but interrupts may cause problems as they preempt running tasks
- *buffer overflow* of the event queue may happen if too many events are generated by the environment or by tasks (guarantee requires bounded behavior of environment)
- *tasks with a long running time* prevent other tasks from running and may cause buffer overflow as no events are being processed during this time
 - partition tasks into smaller ones
 - but the local context must be stored




Preemptive Event-Triggered Scheduling – Stack Policy

- Each event/task has a *fixed priority*. Tasks with *higher priority can preempt* tasks with lower priority. This partly solves the problem of long-running tasks.
- If *the order of preemption is restricted*, we can use the usual stack-based context mechanism of function calls. The context of a function contains the necessary state such as local variables and saved registers.



Preemptive Event-Triggered Scheduling – Stack Policy



- *Tasks must finish in LIFO (last in first out) order* of their instantiation, that is, the preempting task must finish before the preempted task can continue.
 - this restricts flexibility of the approach
 - not useful if tasks wait some unknown time for external events (i.e., they are blocked)
- *Shared resources* (communication between tasks!) *must be protected*, for example, by disabling interrupts or by the use of semaphores.

Preemptive Event-Triggered Scheduling – Stack Policy

main:

```
while (true) {  
    if (event queue is empty) {  
        sleep();  
    } else {  
        select event from event queue;  
        execute selected task;  
        remove selected event from queue;  
    }  
}
```

set CPU to low power mode;
processing continues after interrupt

for example using a function pointer
in C; task returns after finishing.

InsertEvent:

```
put new event into event queue;  
select event from event queue;  
if (selected task  $\neq$  running task) {  
    execute selected task;  
    remove selected event from queue;  
}  
return;
```

Interrupt:

```
InsertEvent (...);  
return;
```

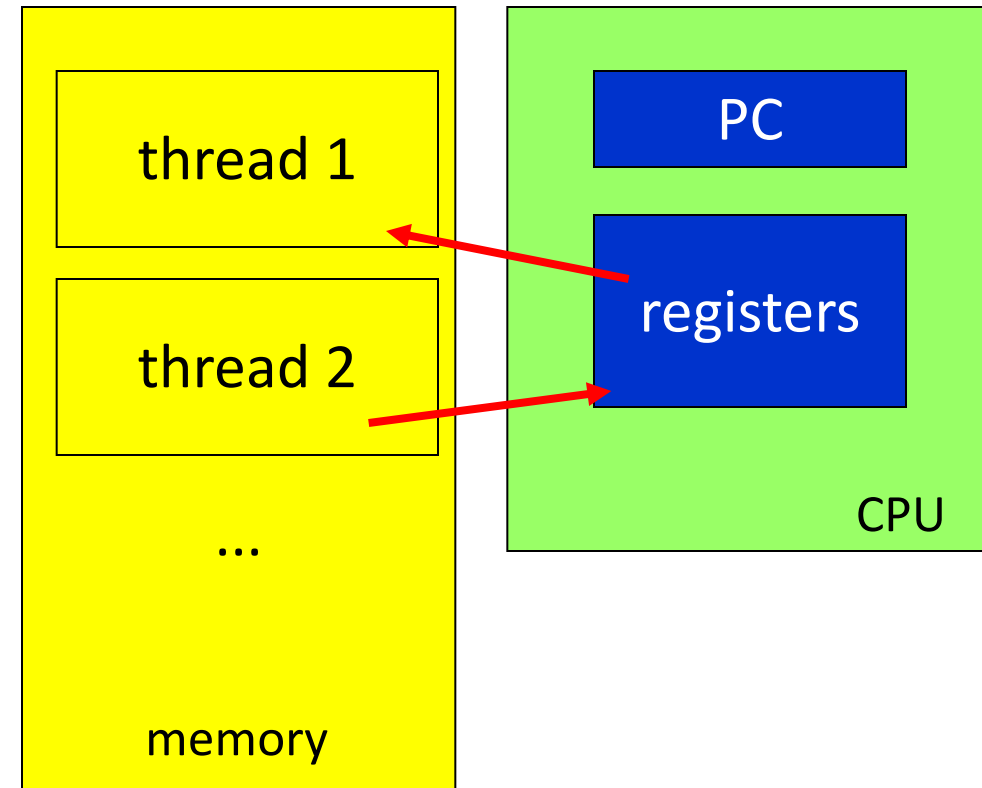
may be called by interrupt service
routines (ISR) or tasks

Thread

- *A thread is a unique execution of a piece of code.*
 - Several copies of such a “program” may run simultaneously or at different times.
 - Threads share the same processor and its peripherals.
- *A thread has its own local state.* This state consists mainly of:
 - register values;
 - memory stack (local variables);
 - program counter;
- *Several threads may have a shared state* consisting of global variables.

Threads and Memory Organization

- *Activation record*, or *thread context*, contains the thread-local state, including registers and local data structures.
- Each thread has a *fixed memory region* for storing its context.
- *Context switch*:
 - current CPU context (program counter, registers) goes out
 - new CPU context goes in



Co-operative Multitasking

- *Each thread allows a context switch to another thread* at a call to the `cswitch()` function.
 - This function is part of the underlying runtime system (operating system).
 - A *scheduler* within this runtime system chooses which thread will run next. This could be a different thread, or the same one in case no other thread is ready to run.
- ***Advantages:***
 - predictable, where context switches can occur (programmer has full control)
 - less errors with use of shared resources if the switch locations are chosen carefully
- ***Disadvantages:***
 - programming errors (e.g., if the `cswitch()` function is never called) can keep other threads out as the running thread may never give up the CPU
 - real-time behavior may be at risk if a thread runs for too long before the next context switch is allowed

Example: Co-operative Multitasking

Thread 1

```
if (x > 2)
    sub1(y);
else
    sub2(y);
cswitch();
proca(a,b,c);
```

Thread 2

```
procd(r,s,t);
cswitch();
if (val1 == 3)
    abc(val2);
rst(val3);
```

Scheduler

```
save_state(current);
p = choose_process();
load_and_go(p);
```

Preemptive Multitasking

- *Most general form of multitasking:*

- The scheduler in the runtime system (operating system) controls when contexts switches take place.
- The scheduler also determines what thread runs next.

- *State diagram corresponding to each single thread:*

- **Run:** A thread enters this state as it starts executing on the processor. Only one thread can be in this state.
- **Ready:** State of threads that are ready to execute but cannot be executed because the processor is assigned to another thread.
- **Blocked:** A task enters this state when it waits for an event.

