

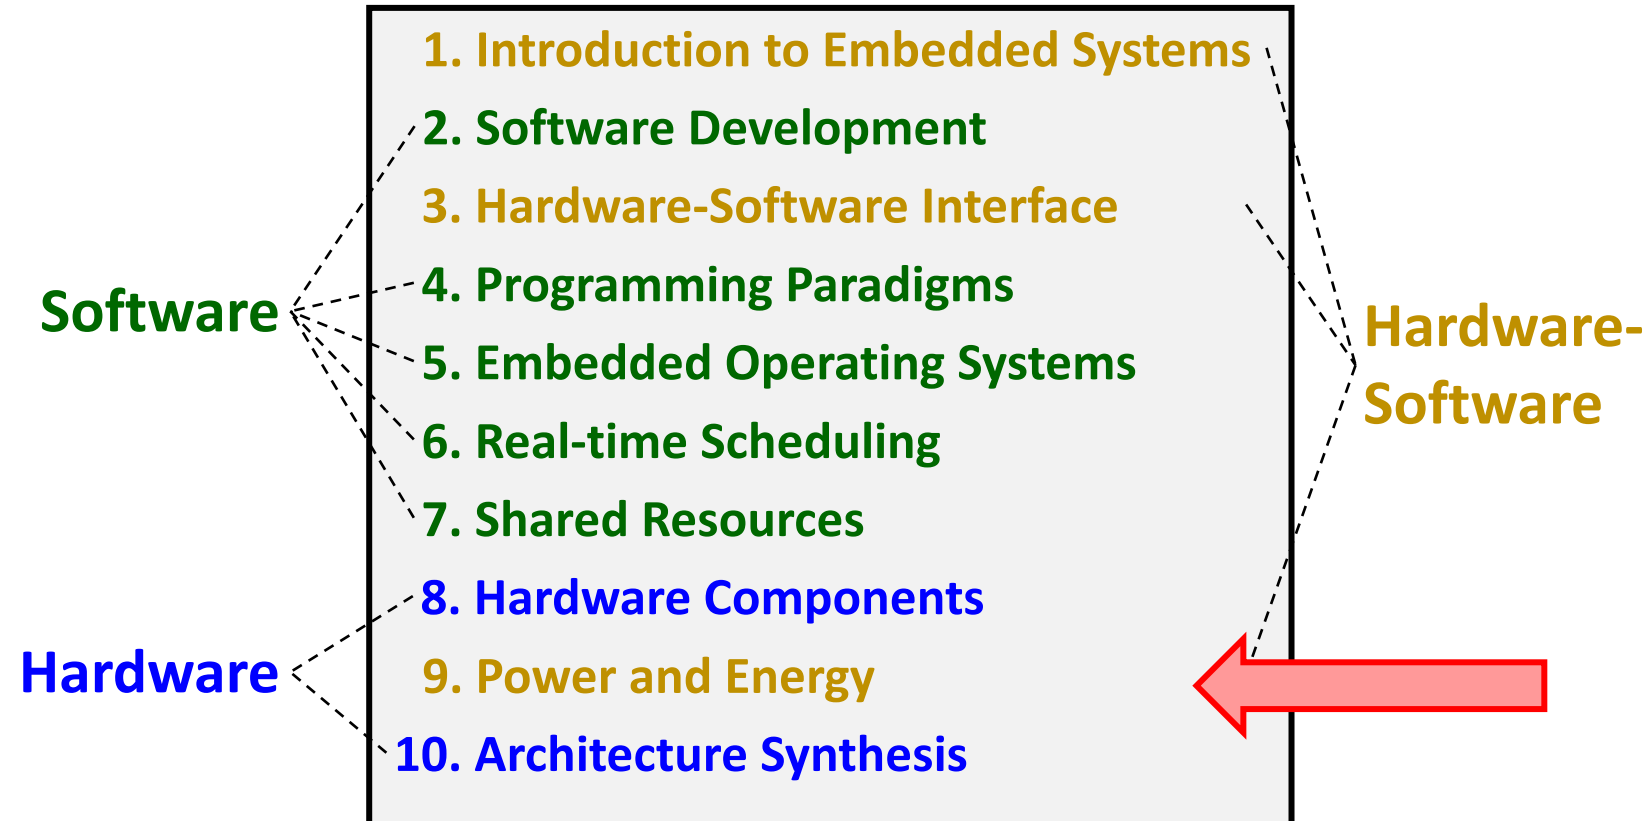
Introduction to Embedded Systems

9. Power and Energy

Prof. Dr. Marco Zimmerling



Where we are ...



General Remarks

Power and Energy Consumption

- Statements that are true since a decade or longer:


„Power is considered as the most important constraint in embedded systems.“ [in: L. Eggermont (ed): Embedded Systems Roadmap 2002, STW]

“Power demands are increasing rapidly, yet battery capacity cannot keep up.” [in Ditzel et al.: Power-Aware Architecting for data-dominated applications, 2007, Springer]

- *Main reasons* are:
 - power provisioning is expensive
 - battery capacity is growing only slowly
 - devices may overheat
 - energy harvesting (e.g., from solar cells) is limited due to low energy density

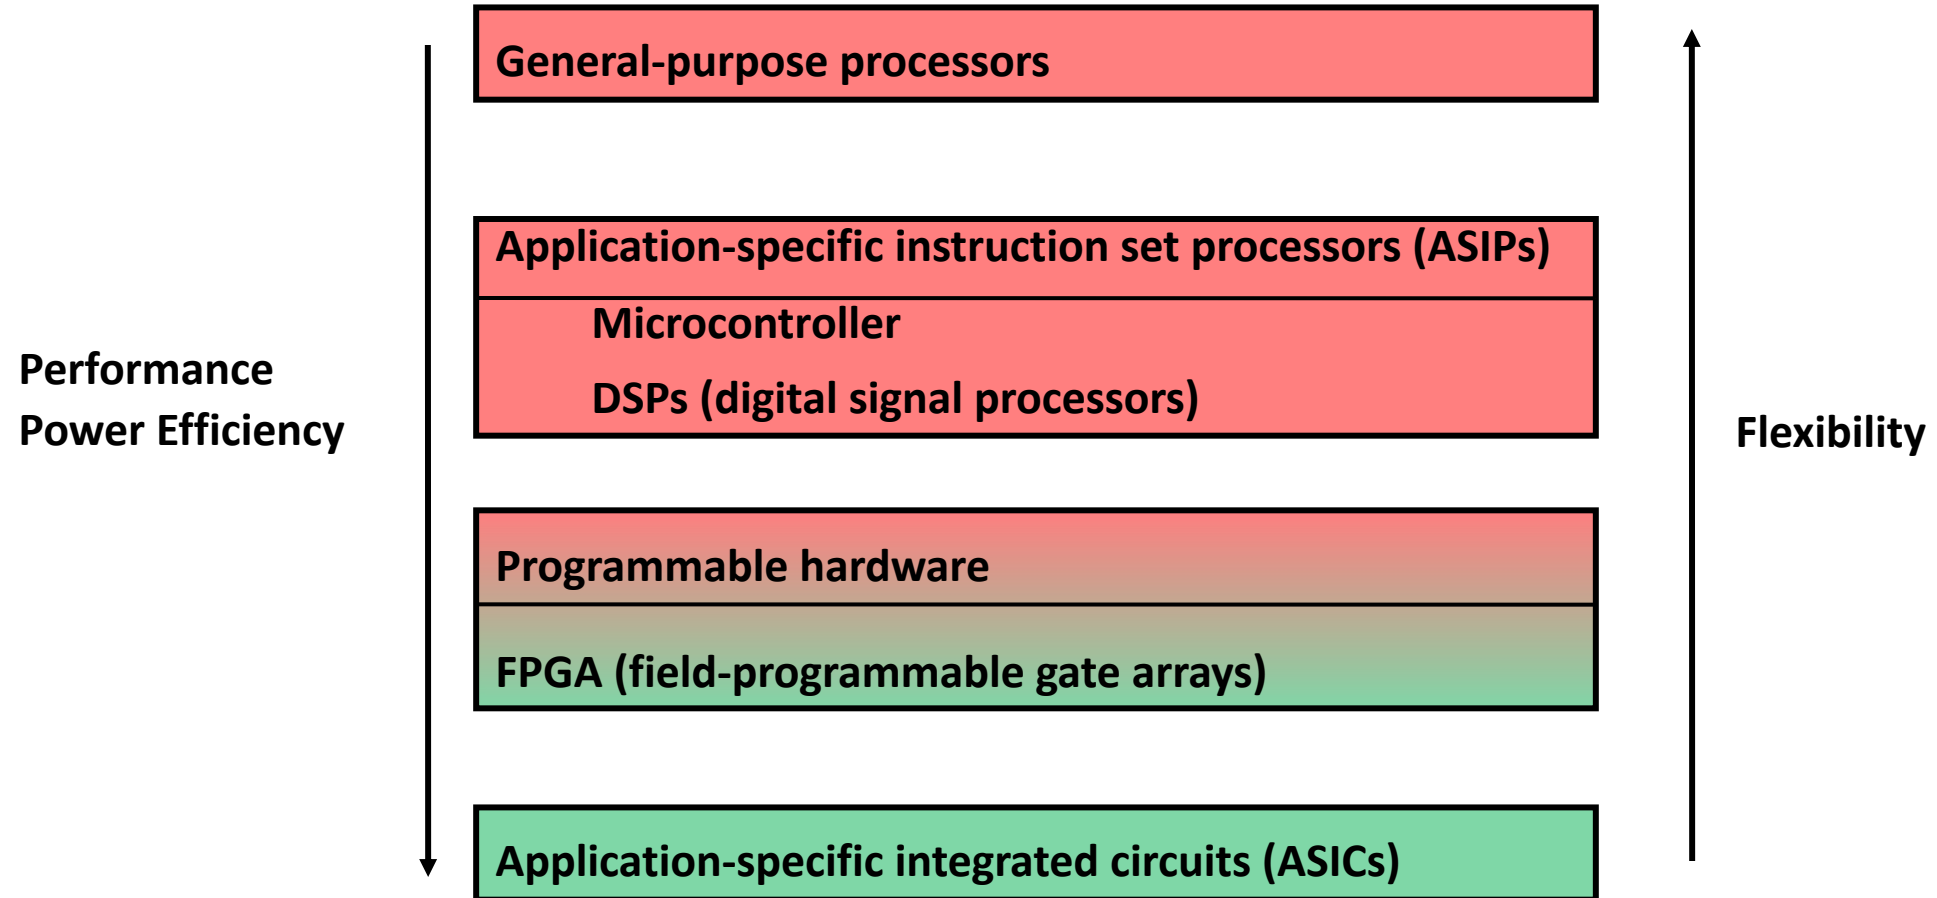


Some Trends




Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2017 by K. Rupp

Implementation Alternatives



Energy Efficiency

- It is necessary to *optimize HW and SW*.
- Use *heterogeneous architectures* in order to adapt to required performance and application.
- Apply *specialization techniques*:
 - Higher parallelism
 - Turn off unused components (e.g., low-power modes)
 - Higher heterogeneity
 - Voltage and frequency scaling




Power and Energy

Power and Energy

$$E = \int P(t)dt = \int \underset{\text{voltage}}{V(t)} \cdot \underset{\text{current}}{I(t)}dt$$


Power is an *instantaneous quantity* measured at time t



Energy is consumed *by a given task* over a certain time interval


Because energy is not an instantaneous quantity, it is important to specify for *what* a given amount of energy is used (context)!

Power and Energy




In some cases, faster execution also means less energy, but the opposite may be true if power has to be increased to allow for a faster execution.

Power and Energy




In some cases, faster execution also means less energy, but the opposite may be true if power has to be increased to allow for a faster execution.

Power and Energy



In some cases, faster execution also means less energy, but the opposite may be true if power has to be increased to allow for a faster execution.

Power and Energy



Common case: Power must be increased to execute faster.

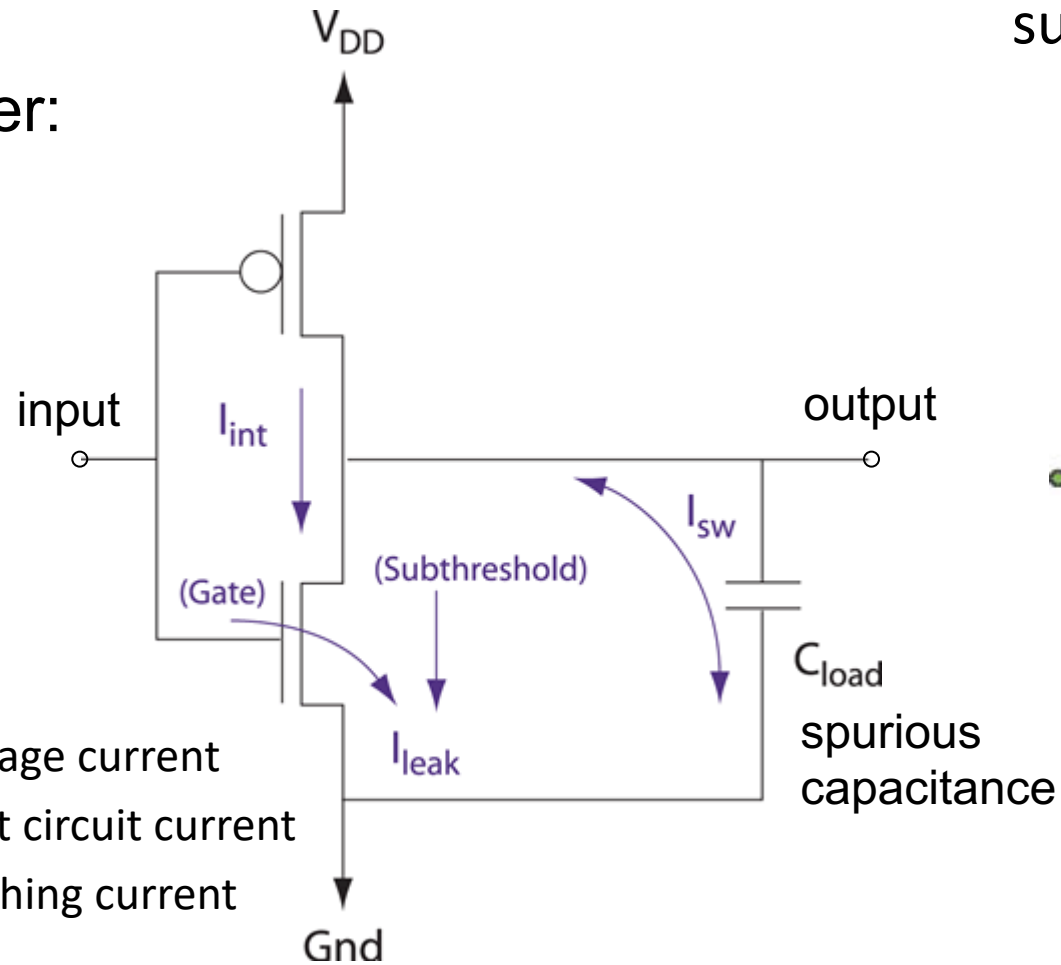
In some cases, faster execution also means less energy, but the opposite may be true if power has to be increased to allow for a faster execution.

Low Power vs. Low Energy


- Power and energy impact system design in different ways
- Minimizing the *power consumption* (= *voltage x current*) is important for
 - the design of the power supply and voltage regulators
 - the dimensioning of interconnect between power supply and components (e.g., reduced power enables the use of thinner wires)
 - cooling (short term cooling)
 - high cost, limited space
- Minimizing the *energy consumption* is important due to
 - restricted availability of energy (e.g., in mobile systems or IoT devices)
 - limited battery capacities (only slowly improving)
 - very high costs of energy (energy harvesting, solar panels, maintenance/batteries)
 - long lifetimes, low temperatures

Power Consumption of a CMOS Gate

Inverter:




subthreshold (I_{SUB}), junction (I_{JUNC}) and gate-oxide (I_{GATE}) leakage



I_{leak} : leakage current
 I_{int} : short circuit current
 I_{sw} : switching current

Main Sources of Power Consumption of a CMOS Processors

- *Dynamic* power consumption:
 - (Dis)charging capacitances while switching
 - Short-circuit power consumption due to a short-circuit path between supply rails while switching
- *Static* power consumption:
 - Gate-oxide, subthreshold, and junction leakage while nothing happens (i.e., no clock, no switching of gate inputs)



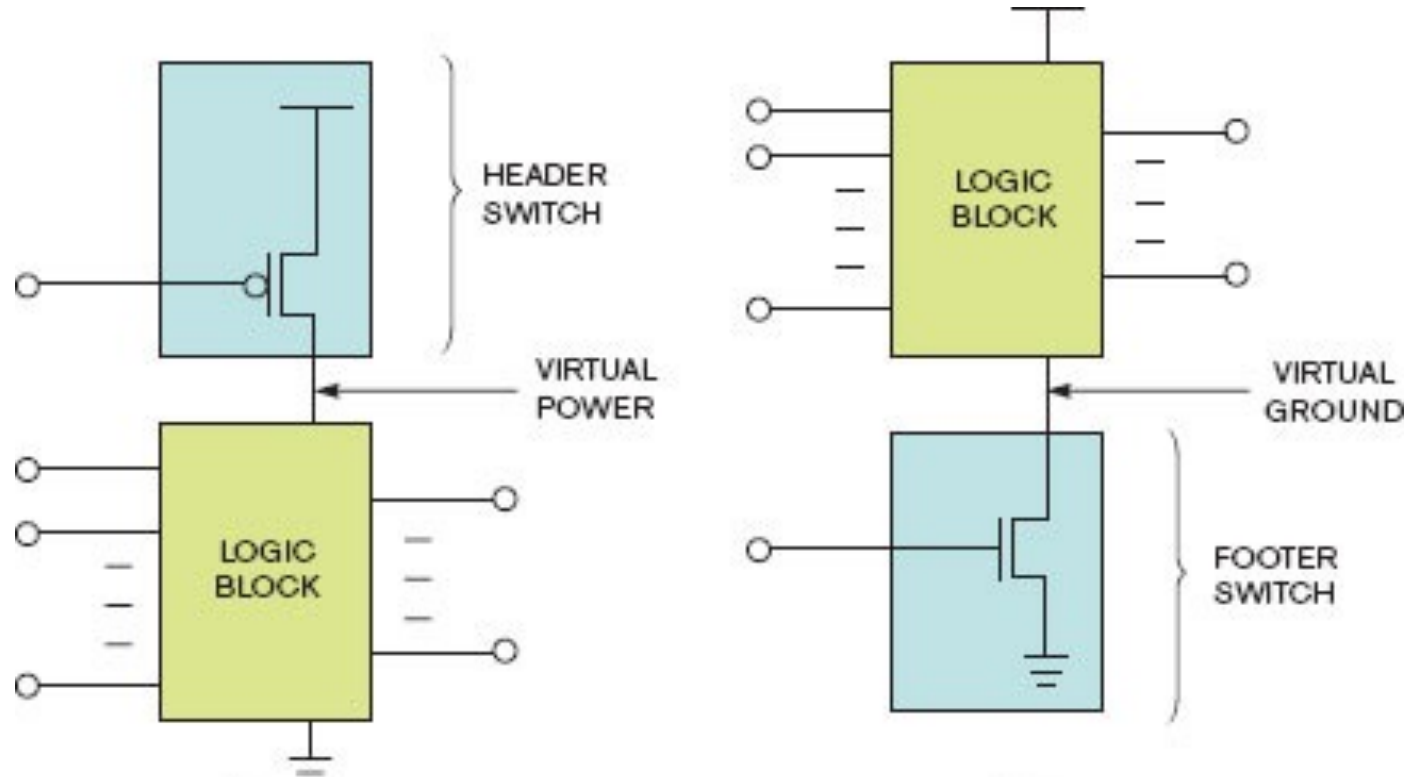
[J. Xue, T. Li, Y. Deng, Z. Yu, Full-chip leakage analysis for 65 nm CMOS technology and beyond, Integration VLSI J. 43 (4) (2010) 353–364]

Techniques to Reduce Static Power

Power Supply Gating

Power gating is one of the most effective ways to reduce static power consumption (leakage)

- Idea: Cut power supply off when components are unused (duty cycling, low-power modes)



Techniques to Reduce Dynamic Power

Voltage Scaling: A Simple Analytical Model

Average power consumption of CMOS circuits (ignoring leakage):

$$P \sim \alpha C_L V_{dd}^2 f$$

V_{dd} : supply voltage
 α : switching activity
 C_L : load capacity
 f : clock frequency

Delay of CMOS circuits:

$$\tau \sim C_L \frac{V_{dd}}{(V_{dd} - V_T)^2} \sim \frac{C_L}{V_{dd}}$$

V_{dd} : supply voltage
 V_T : threshold voltage
 $V_T \ll V_{dd}$

Decreasing V_{dd} reduces P quadratically (assuming f is constant).

But decreasing V_{dd} also increases the gate delay τ reciprocally.

Maximal frequency f_{\max} decreases linearly with decreasing V_{dd} (i.e., “reducing the voltage makes the system slower”).

$$f_{\max} \sim \frac{1}{\tau} \sim \frac{V_{dd}}{C_L}$$

Voltage Scaling: A Simple Analytical Model

$$P \sim \alpha C_L V_{dd}^2 f$$

This is the energy per cycle.


$$E \sim \alpha C_L V_{dd}^2 f t = \overbrace{\alpha C_L V_{dd}^2} (\#cycles)$$

This is the energy needed for executing a task that requires #cycles (i.e., number of cycles).

Saving energy for a given task:

- reduce the supply voltage V_{dd}
- reduce switching activity α
- reduce the load capacitance C_L
- reduce the number of cycles $\#cycles$


Parallelism



$$E \sim V_{dd}^2 (\text{\#cycles})$$

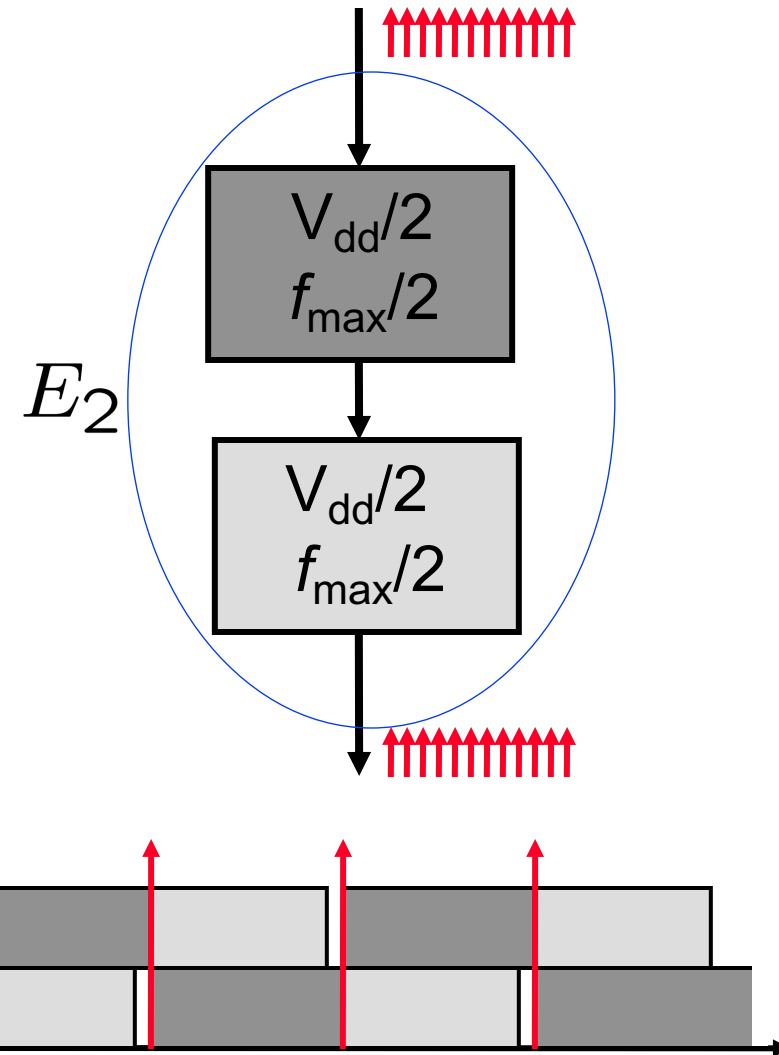
$$E_2 = \frac{1}{4} E_1$$

Pipelining



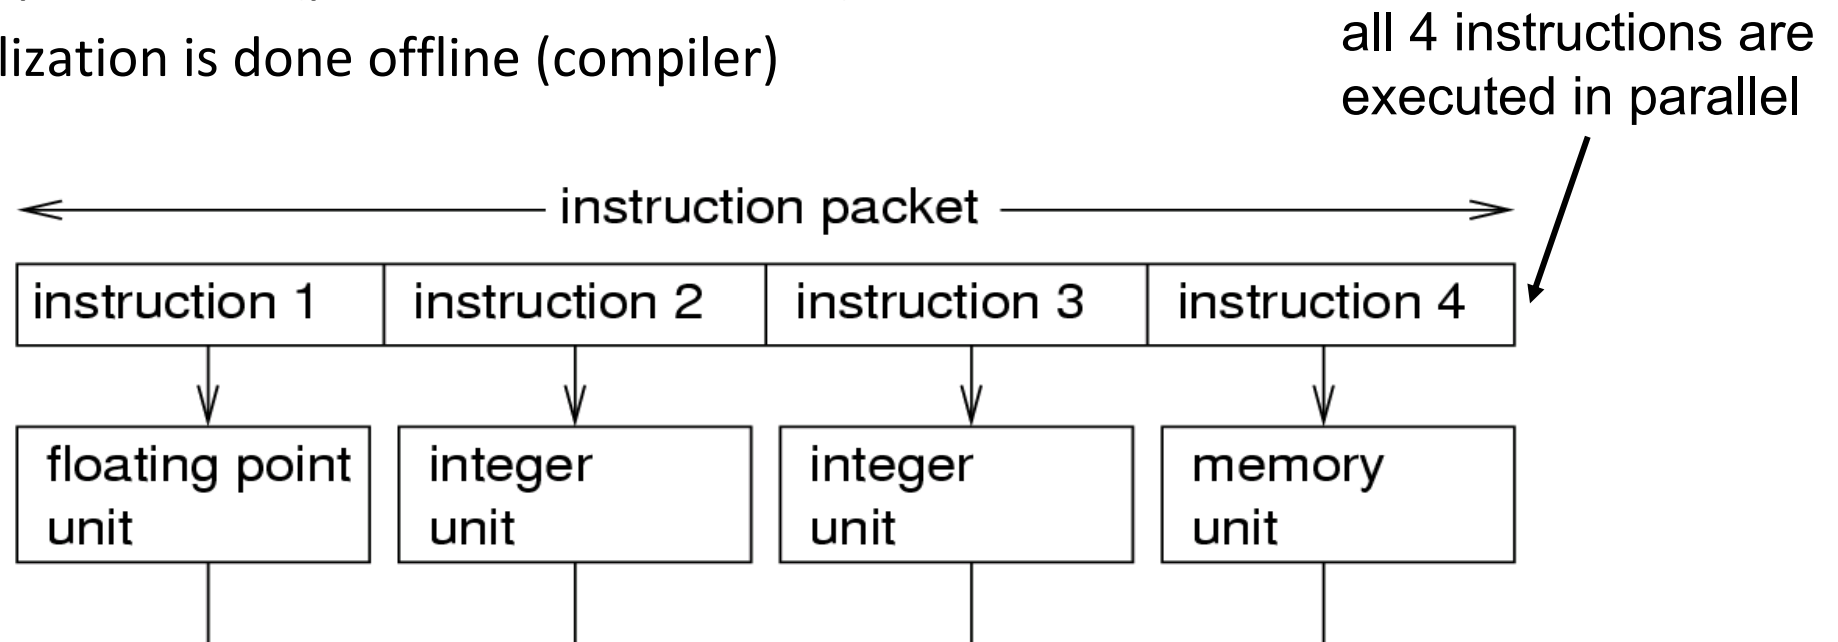
$$E \sim V_{dd}^2 (\text{\#cycles})$$

$$E_2 = \frac{1}{4} E_1$$



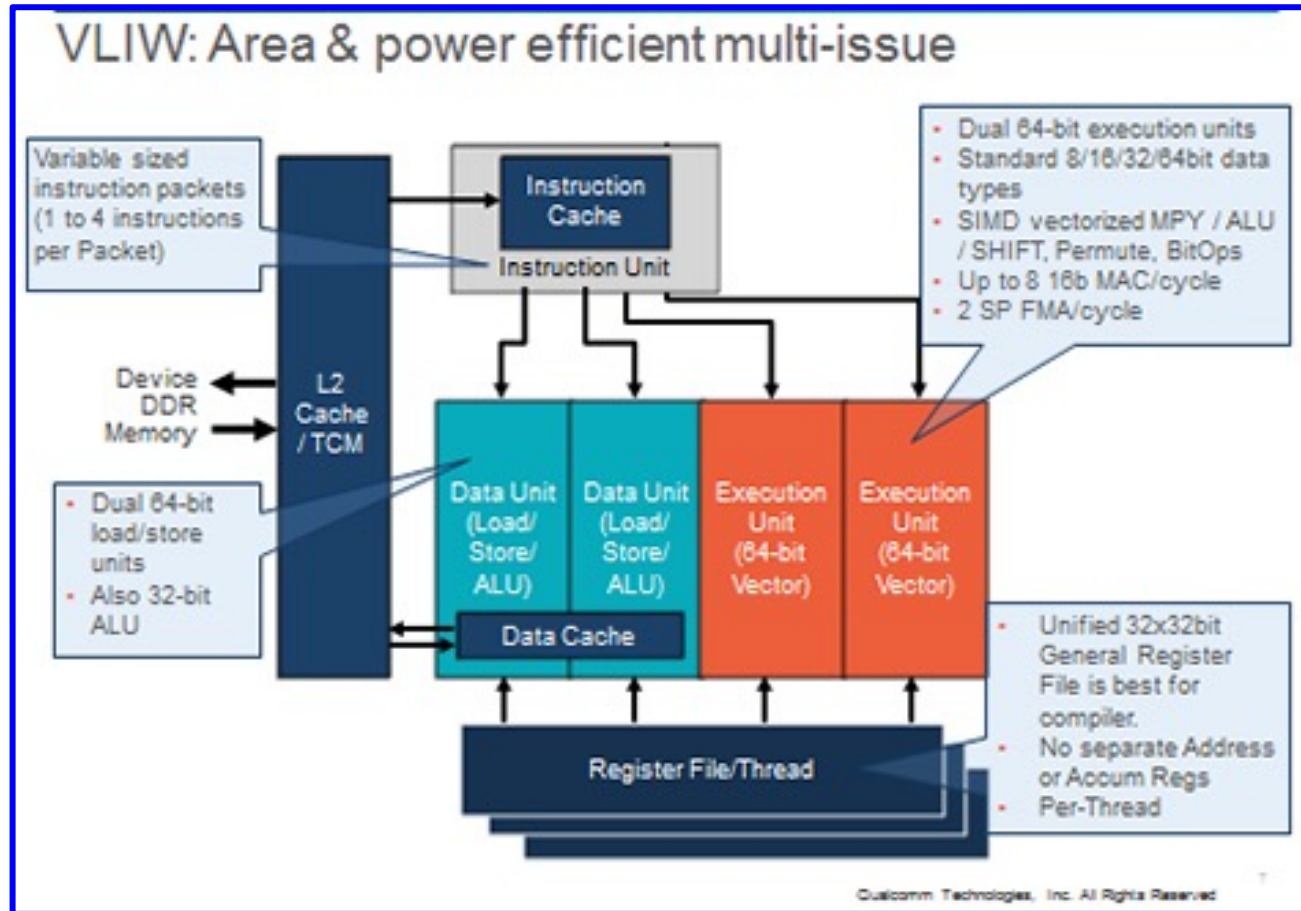
VLIW (Very Long Instruction Word) Architectures

- **Large degree of parallelism**
 - many parallel computational units, (deeply) pipelined
- **Simple hardware architecture**
 - explicit parallelism (parallel instruction set)
 - parallelization is done offline (compiler)

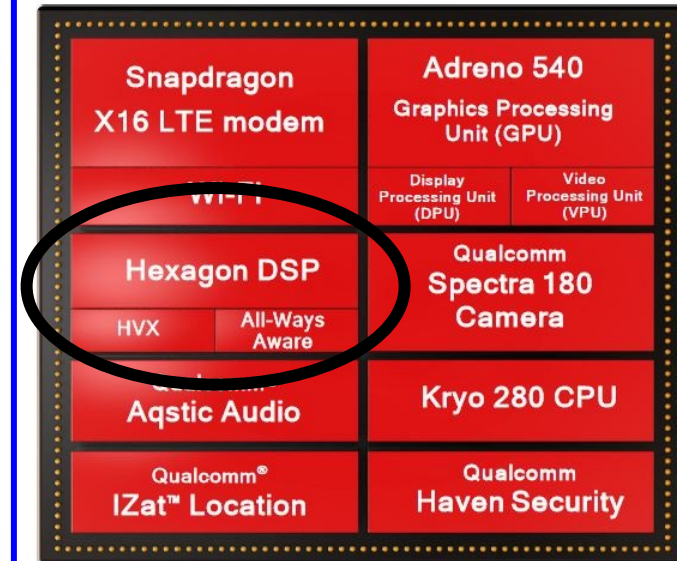


Example: Qualcomm Hexagon

Hexagon DSP



Snapdragon 835 (Galaxy S8)



So far: *Statically* decrease voltage and frequency

Next: *Dynamically* change voltage and frequency

Dynamic Voltage and Frequency Scaling (DVFS)

$$P \sim \alpha C_L V_{dd}^2 f$$

energy per cycle reduce voltage -> reduce energy per task

$$E \sim \alpha C_L V_{dd}^2 f t = \alpha C_L V_{dd}^2 (\#cycles)$$

$$f \sim \frac{1}{\tau} \sim V_{dd}$$

reduce voltage -> reduce clock frequency


maximum
frequency
of operation

gate delay

Saving energy for a given task:

- reduce the supply voltage V_{dd}
- reduce switching activity α
- reduce the load capacitance C_L
- reduce the number of cycles $\#cycles$


Example: Dynamic Voltage and Frequency Scaling



Example: DVFS – Complete Task as Early as Possible

| V_{dd} [V] | 5.0 | 4.0 | 2.5 |
|-----------------------|-----|-----|-----|
| Energy per cycle [nJ] | 40 | 25 | 10 |
| f_{max} [MHz] | 50 | 40 | 25 |
| cycle time [ns] | 20 | 25 | 40 |

We suppose a task that needs 10^9 cycles to execute within 25 seconds.




$$E_a = 10^9 \times 40 \times 10^{-9} = 40 \text{ [J]}$$

Example: DVFS – Use Two Voltages

| V_{dd} [V] | 5.0 | 4.0 | 2.5 |
|-----------------------|-----|-----|-----|
| Energy per cycle [nJ] | 40 | 25 | 10 |
| f_{max} [MHz] | 50 | 40 | 25 |
| cycle time [ns] | 20 | 25 | 40 |


b) [V²] 750M cycles @ 50 MHz + 250M cycles @ 25 MHz



$$\begin{aligned} E_b &= 750 \cdot 10^6 \times 40 \times 10^{-9} \\ &\quad + 250 \cdot 10^6 \times 10 \times 10^{-9} \\ &= 32.5 \text{ [J]} \end{aligned}$$


Example: DVFS – Use One Voltage

| V_{dd} [V] | 5.0 | 4.0 | 2.5 |
|-----------------------|-----|-----|-----|
| Energy per cycle [nJ] | 40 | 25 | 10 |
| f_{max} [MHz] | 50 | 40 | 25 |
| cycle time [ns] | 20 | 25 | 40 |



$$E_c = 10^9 \times 25 \times 10^{-9} \\ = 25 \text{ [J]}$$

DVFS: Optimal Strategy



Execute task in fixed time T
with variable voltage $V_{dd}(t)$:


$$\text{gate delay: } \tau \sim \frac{1}{V_{dd}}$$

$$\text{execution rate: } f(t) \sim V_{dd}(t)$$

$$\text{invariant: } \int V_{dd}(t) dt = \text{const.}$$

- **case A:** execute at voltage x for $T \cdot a$ time units and at voltage y for $(1-a) \cdot T$ time units;
energy consumption: $T \cdot (P(x) \cdot a + P(y) \cdot (1-a))$

DVFS: Optimal Strategy



Execute task in fixed time T
with variable voltage $V_{dd}(t)$:

gate delay: $\tau \sim \frac{1}{V_{dd}}$

execution rate: $f(t) \sim V_{dd}(t)$

invariant: $\int V_{dd}(t) dt = \text{const.}$

- **case A**: execute at voltage x for $T \cdot a$ time units and at voltage y for $(1-a) \cdot T$ time units;


energy consumption: $T \cdot (P(x) \cdot a + P(y) \cdot (1-a))$

- **case B**: execute at voltage $z = a \cdot x + (1-a) \cdot y$ for T time units;

energy consumption: $T \cdot P(z)$

Ensures that we compare
the energies consumed by
executing the same algo-
rithm with the same #cycles

DVFS: Optimal Strategy



$$z \cdot T = a \cdot T \cdot x + (1-a) \cdot T \cdot y$$

$$z = a \cdot x + (1-a) \cdot y$$

Execute task in fixed time T with variable voltage $V_{dd}(t)$:

gate delay: $\tau \sim \frac{1}{V_{dd}}$

execution rate: $f(t) \sim V_{dd}(t)$

invariant: $\int V_{dd}(t) dt = \text{const.}$

Ensures that we compare the energies consumed by executing the same algorithm with the same #cycles


- **case A:** execute at voltage x for $T \cdot a$ time units and at voltage y for $(1-a) \cdot T$ time units;

energy consumption: $T \cdot (P(x) \cdot a + P(y) \cdot (1-a))$

- **case B:** execute at voltage $z = a \cdot x + (1-a) \cdot y$ for T time units;

energy consumption: $T \cdot P(z)$

DVFS: Optimal Strategy



If possible, running at a constant frequency (voltage) minimizes the energy consumption for dynamic voltage scaling:

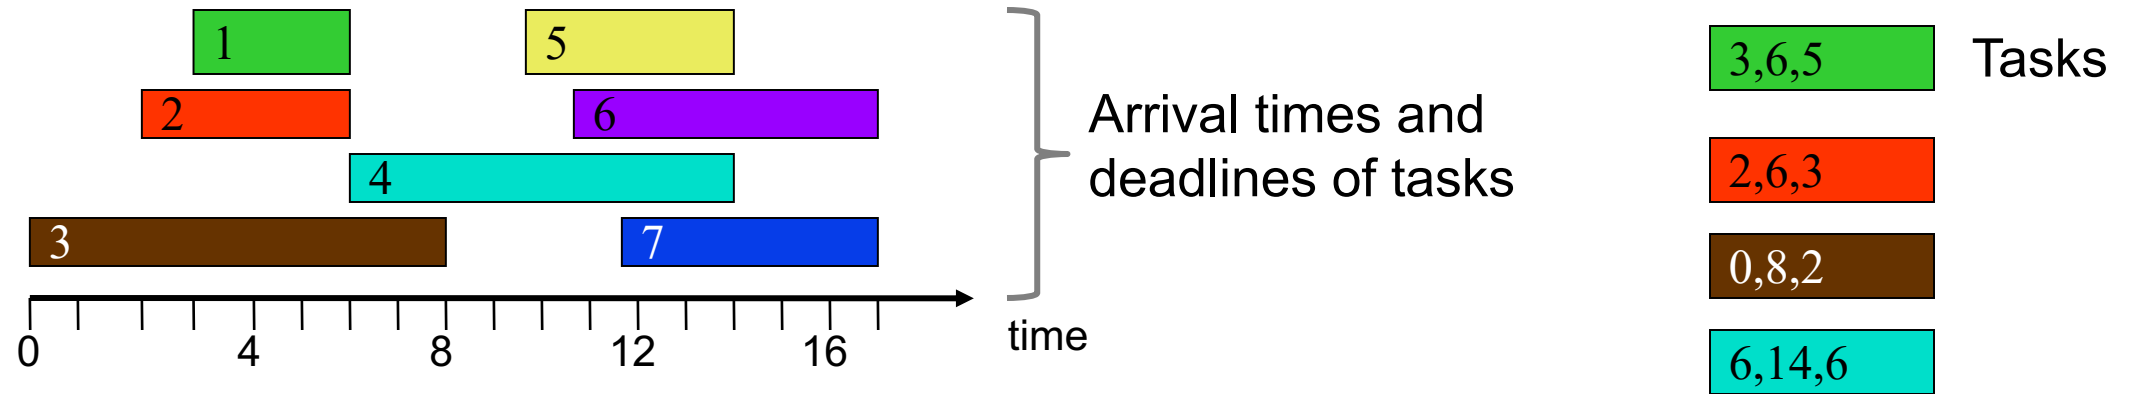
case A is always worse if the power consumption is a convex function of the supply voltage

DVFS: Real-Time Offline Scheduling on One Processor

- Let us model a set of independent tasks as follows:
 - We suppose that a task $\tau_i \in V$
 - requires c_i computation time at normalized processor frequency 1
 - arrives at time a_i
 - has (absolute) deadline constraint d_i
- How do we schedule these tasks such that all these tasks can be finished **no later than their deadlines** and the energy consumption is **minimized**?
 - YDS Algorithm from “A Scheduling Model for Reduced CPU Energy”, Frances Yao, Alan Demers, and Scott Shenker, FOCS 1995.”

If possible, running at a constant frequency (voltage) minimizes the energy consumption for dynamic voltage scaling.

YDS Optimal DVFS Algorithm for Offline Scheduling



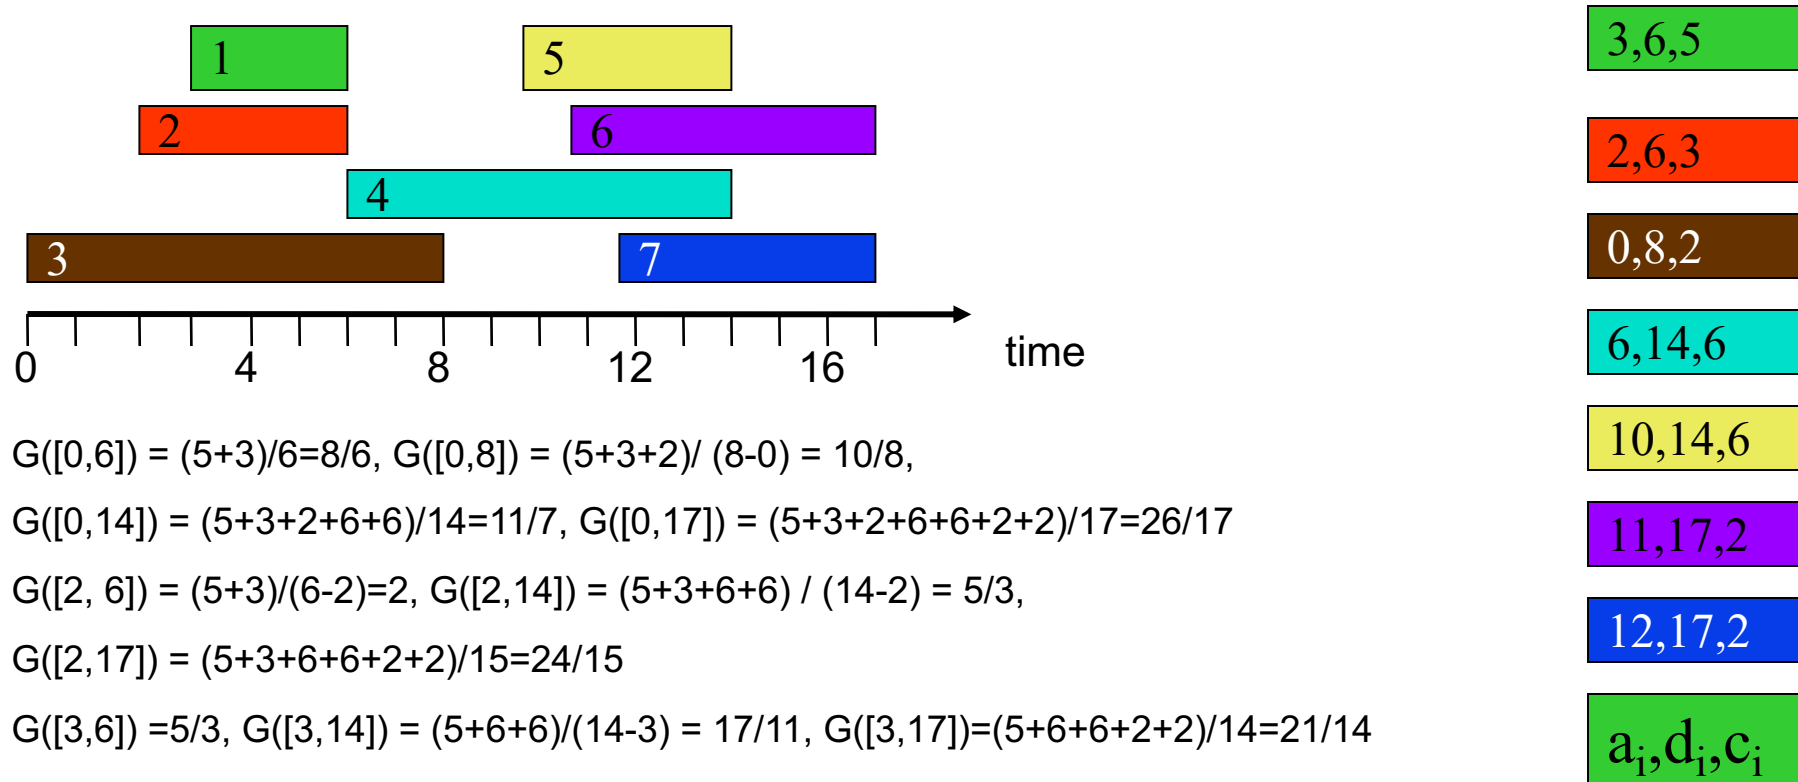
- Define **intensity** $G([z, z'])$ in some time interval $[z, z']$:
 - average accumulated execution time of all tasks that have arrival and deadline in $[z, z']$ relative to the length of the interval $z'-z$

$$V'([z, z']) = \{v_i \in V : z \leq a_i < d_i \leq z'\}$$

$$G([z, z']) = \sum_{v_i \in V'([z, z'])} c_i / (z' - z)$$


YDS Optimal DVFS Algorithm for Offline Scheduling

Step 1: Execute jobs in the interval with the highest intensity by using the earliest-deadline first schedule and running at the intensity as the frequency.



YDS Optimal DVFS Algorithm for Offline Scheduling

Step 1: Execute jobs in the interval with the highest intensity by using the earliest-deadline first schedule and running at the intensity as the frequency.



$$G([0,6]) = (5+3)/6=8/6, G([0,8]) = (5+3+2)/(8-0) = 10/8,$$

$$G([0,14]) = (5+3+2+6+6)/14=11/7, G([0,17]) = (5+3+2+6+6+2+2)/17=26/17$$

*critical
interval*

$$G([2,6]) = (5+3)/(6-2)=2 \quad G([2,14]) = (5+3+6+6)/(14-2) = 5/3,$$

$$G([2,17]) = (5+3+6+6+2+2)/15=24/15$$

$$G([3,6]) = 5/3, G([3,14]) = (5+6+6)/(14-3) = 17/11, G([3,17])=(5+6+6+2+2)/14=21/14$$

$$G([6,14]) = 12/(14-6)=12/8, G([6,17]) = (6+6+2+2)/(17-6)=16/11$$

$$G([10,14]) = 6/4, G([10,17]) = 10/7, G([11,17]) = 4/6, G([12,17]) = 2/5$$

3,6,5

2,6,3

0,8,2

6,14,6

10,14,6


11,17,2

12,17,2

a_i, d_i, c_i


YDS Optimal DVFS Algorithm for Offline Scheduling

Step 1: Execute jobs in the interval with the highest intensity by using the earliest-deadline first schedule and running at the intensity as the frequency.



YDS Optimal DVFS Algorithm for Offline Scheduling

Step 2: Adjust the arrival times and deadlines by excluding the possibility to execute within the previous critical intervals.



0,8,2

6,14,6

10,14,6



11,17,2

12,17,2

a_i, d_i, c_i


YDS Optimal DVFS Algorithm for Offline Scheduling

Step 2: Adjust the arrival times and deadlines by excluding the possibility to execute within the previous critical intervals.



YDS Optimal DVFS Algorithm for Offline Scheduling

Step 3: Run the algorithm for the revised input again



$$G([0,4])=2/4, G([0,10]) = 14/10, G([0,13])=18/13$$

$$G([2,10])=12/8, G([2,13]) = 16/11, G([6,10])=6/4$$

$$G([6,13])=10/7, G([7,13])=4/6, G([8,13])=4/5$$

0,4,2

2,10,6

6,10,6


7,13,2

8,13,2

a_i, d_i, c_i

YDS Optimal DVFS Algorithm for Offline Scheduling

Step 3: Run the algorithm for the revised input again



$$G([0,4])=2/4, G([0,10]) = 14/10, G([0,13])=18/13$$

$$G([2,10])=12/8, G([2,13]) = 16/11, G([6,10])=6/4$$

$$G([6,13])=10/7, G([7,13])=4/6, G([8,13])=4/5$$

0,4,2

2,10,6

6,10,6


7,13,2

8,13,2

a_i, d_i, c_i

YDS Optimal DVFS Algorithm for Offline Scheduling


Step 3: Run the algorithm for the revised input again



$$G([0,4])=2/4, G([0,10]) = 14/10, G([0,13])=18/13$$

$$G([2,10])=12/8, G([2,13]) = 16/11, G([6,10])=6/4$$

$$G([6,13])=10/7, G([7,13])=4/6, G([8,13])=4/5$$



0,4,2

2,10,6

6,10,6

7,13,2




8,13,2

a_i, d_i, c_i

YDS Optimal DVFS Algorithm for Offline Scheduling

Step 3: Run the algorithm for the revised input again


Step 4: Put pieces together



energy-
optimal
schedule

| task | v_1 | v_2 | v_3 | v_4 | v_5 | v_6 | v_7 |
|-----------|-------|-------|-------|-------|-------|-------|-------|
| frequency | 2 | 2 | 1 | 1.5 | 1.5 | 4/3 | 4/3 |

YDS Optimal DVFS Algorithm for Online Scheduling




0,8,2

Continuously update to the best schedule for all arrived tasks:

Time 0: task v_3 is executed at 2/8

a_i, d_i, c_i

YDS Optimal DVFS Algorithm for Online Scheduling



2,6,3

0,8,2

Continuously update to the best schedule for all arrived tasks:


Time 0: task v_3 is executed at $2/8$

Time 2: task v_2 arrives

- $G([2,6]) = 3/4$, $G([2,8]) = 4.5/6 = 3/4$ \Rightarrow execute v_3, v_2 at $3/4$

a_i, d_i, c_i

YDS Optimal DVFS Algorithm for Online Scheduling



3,6,5

2,6,3

0,8,2

Continuously update to the best schedule for all arrived tasks:

Time 0: task v_3 is executed at $2/8$

Time 2: task v_2 arrives


- $G([2,6]) = \frac{3}{4}$, $G([2,8]) = \frac{4.5}{6} = \frac{3}{4}$ \Rightarrow execute v_3, v_2 at $\frac{3}{4}$

Time 3: task v_1 arrives

- $G([3,6]) = \frac{(5+3-\frac{3}{4})}{3} = \frac{29}{12}$, $G([3,8]) < G([3,6]) \Rightarrow$ execute v_2 and v_1 at $\frac{29}{12}$

a_i, d_i, c_i

YDS Optimal DVFS Algorithm for Online Scheduling



Continuously update to the best schedule for all arrived tasks:

Time 0: task v_3 is executed at $2/8$

Time 2: task v_2 arrives

- $G([2,6]) = \frac{3}{4}$, $G([2,8]) = \frac{4.5}{6} = \frac{3}{4} \Rightarrow$ execute v_3, v_2 at $\frac{3}{4}$

Time 3: task v_1 arrives


- $G([3,6]) = \frac{(5+3-3/4)}{3} = \frac{29}{12}$, $G([3,8]) < G([3,6]) \Rightarrow$ execute v_2 and v_1 at $\frac{29}{12}$

Time 6: task v_4 arrives

- $G([6,8]) = \frac{1.5}{2}$, $G([6,14]) = \frac{7.5}{8} \Rightarrow$ execute v_3 and v_4 at $\frac{15}{16}$

a_i, d_i, c_i

YDS Optimal DVFS Algorithm for Online Scheduling



3,6,5

2,6,3

0,8,2

6,14,6

10,14,6

a_i, d_i, c_i

Continuously update to the best schedule for all arrived tasks:

Time 0: task v_3 is executed at $2/8$

Time 2: task v_2 arrives

- $G([2,6]) = \frac{3}{4}$, $G([2,8]) = \frac{4.5}{6} = \frac{3}{4} \Rightarrow$ execute v_8, v_2 at $\frac{3}{4}$

Time 3: task v_1 arrives

- $G([3,6]) = \frac{(5+3-\frac{3}{4})}{3} = \frac{29}{12}$, $G([3,8]) < G([3,6]) \Rightarrow$ execute v_2 and v_1 at $\frac{29}{12}$


Time 6: task v_4 arrives

- $G([6,8]) = \frac{1.5}{2}$, $G([6,14]) = \frac{7.5}{8} \Rightarrow$ execute v_3 and v_4 at $\frac{15}{16}$

Time 10: task v_5 arrives

- $G([10,14]) = \frac{39}{16} \Rightarrow$ execute v_4 and v_5 at $\frac{39}{16}$

YDS Optimal DVFS Algorithm for Online Scheduling



3,6,5

2,6,3

0,8,2

6,14,6

10,14,6

11,17,2

12,17,2

a_i, d_i, c_i

Continuously update to the best schedule for all arrived tasks:

Time 0: task v_3 is executed at $2/8$

Time 2: task v_2 arrives

- $G([2,6]) = 3/4$, $G([2,8]) = 4.5/6 = 3/4 \Rightarrow$ execute v_8, v_2 at $3/4$

Time 3: task v_1 arrives

- $G([3,6]) = (5+3-3/4)/3 = 29/12$, $G([3,8]) < G([3,6]) \Rightarrow$ execute v_2 and v_1 at $29/12$

Time 6: task v_4 arrives

- $G([6,8]) = 1.5/2$, $G([6,14]) = 7.5/8 \Rightarrow$ execute v_3 and v_4 at $15/16$

Time 10: task v_5 arrives

- $G([10,14]) = 39/16 \Rightarrow$ execute v_4 and v_5 at $39/16$

Time 11 and Time 12

- The arrival of v_6 and v_7 does not change the critical interval

Time 14:

- $G([14,17]) = 4/3 \Rightarrow$ execute v_6 and v_7 at $4/3$

Remarks on the YDS Algorithm

■ *Offline*

- The algorithm guarantees the minimal energy consumption while satisfying the timing constraints
- The time complexity is $O(N^3)$, where N is the number of tasks in V
 - Finding the critical interval can be done in $O(N^2)$
 - The number of iterations is at most N
- Exercise:
 - For periodic real-time tasks with deadline=period, running at ***constant speed with 100% utilization*** under EDF has minimum energy consumption while satisfying the timing constraints.

■ *Online*

- Compared to the optimal offline solution, the on-line schedule uses at most 27 times of the minimal energy consumption.